

数据库系统概念

Chapter 1

A **database-management system** (DBMS) is a collection of interrelated **data** and a set of **programs** to access those data.

- The collection of data, usually referred to as the **database**, contains information relevant to an enterprise.
- The primary goal of a DBMS is to provide a way to store and retrieve database information that is both **convenient** and **efficient**.

Purpose of Database Systems

- Keeping organizational information in a file-processing system has a number of major disadvantages:
 - data redundancy and inconsistency
 - difficulty in accessing data
 - data isolation
 - integrity problem
 - atomicity problem
 - concurrent-access anomaly
 - security problem

1. View of Data

- A major purpose of a database system is to provide users with **an abstract view** of the data. That is, the system hides certain details of how the data are stored and maintained.
- **Why need data abstraction?**
To hide the complexity from users through several levels of abstraction and simplify users' interactions with the system.

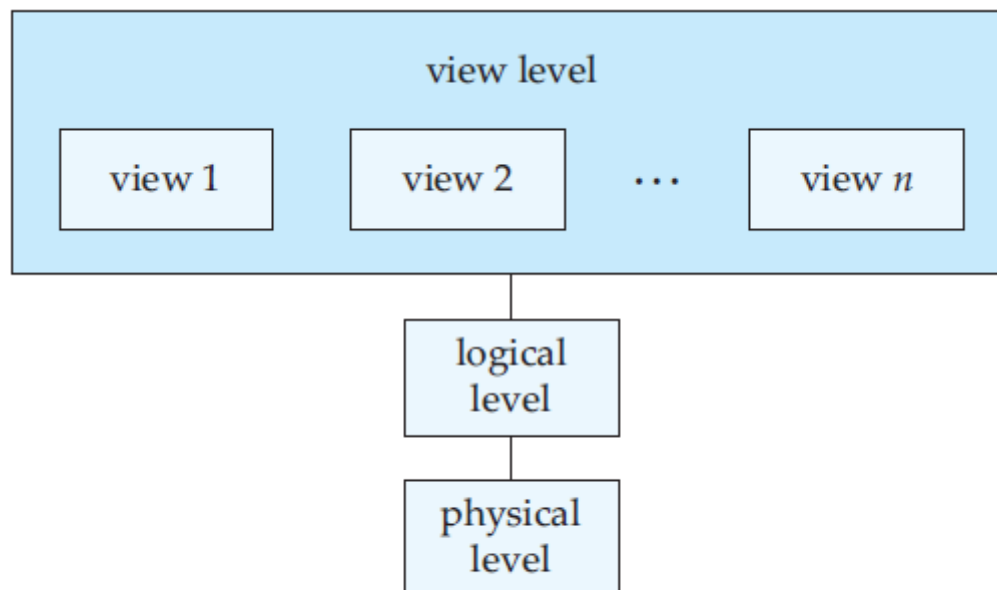


Figure 1.1 The three levels of data abstraction.

- Physical level: The lowest level of abstraction describes *how* the data are actually stored. The physical level describes complex low-level data structures in detail.
- Logical level: The next-higher level of abstraction describes *what* data are stored in the database, and what relationships exist among those data.

Physical data independence: The ability to modify the physical schema without changing the logical schema

Database administrators use the logical level of abstraction.

- View level. The highest level of abstraction describes only part of the entire database.

2. Instances and schemas

The collection of information stored in the database at a particular moment is called an

instance of the database. The overall design of the database is called the database **schema**.

- The **physical schema** describes the database design at the physical level, while the **logical schema** describes the database design at the logical level. A database may also have several schemas at the view level, sometimes called **subschemas**, that describe different views of the database.
- Application programs are said to exhibit **physical data independence** if they do not depend on the physical schema, and thus need not be rewritten if the physical schema changes.

3. Data Models

Underlying the structure of a database is the **data model**.

- Relational Model

- Entity-Relationship Model
- Object-Based Data Model
- Semistructured Data Model

Database Languages

A database system provides a **data-definition language** to specify the database schema and a **data-manipulation language** to express database queries and updates.

- Data-Definition Language, DDL
 - We specify the storage structure and access methods used by the database system by a set of statements in a special type of DDL called a **data storage and definition** language.
 - **Integrity constraints** (*checking when inserting or deleting*):
 - Domain Constraints
 - Referential Integrity
 - Assertions
 - Authorization
 - The DDL gets as input some instructions (statements) and generates some output. The output of the DDL is placed in the **data dictionary**, which contains **metadata**, that is, data about data.
- Data-Manipulation Language, DML

A data-manipulation language (DML) is a language that enables users to access or manipulate data as organized by the appropriate data model.

The types of access are:

- **Retrieval** of information stored in the database
- **Insertion** of new information into the database
- **Deletion** of information from the database
- **Modification** of information stored in the database

There are basically two types of DML:

- Procedural DMLs require a user to specify *what* data are needed and *how* to get those data.
- Declarative DMLs (also referred to as nonprocedural DMLs) require a user to specify *what* data are needed *without* specifying how to get those data.

A **query** is a statement requesting the retrieval of information. The portion of a DML that involves information retrieval is called a **query language**.

Relational Databases

1. Tables

- Each table has multiple columns and each column has a unique name.

2. DML

- The SQL query language is **nonprocedural**. A query takes as input several tables (possibly only one) and always returns **a single table**.

3. DDL

- SQL provides a rich DDL that allows one to define tables, integrity constraints, assertions, etc.

4. Database Access from Application Programs

- To access the database, DML statements need to be executed from the host language. There are two ways to do this:

- By providing an application program interface (set of procedures) that can be used to send DML and DDL statements to the database and retrieve the results.

Like Open Database Connectivity (ODBC) for the *C* language and Java Database Connectivity (JDBC) for the *Java* language.

- By extending the host language syntax to embed DML calls within the host language program.

Usually, a special character prefaces DML calls, and a preprocessor, called the **DML precompiler**, converts the DML statements to normal procedure calls in the host language.

Database Design

- The process of designing the general structure of the database:
 - Logical Design – Deciding on the database schema. Database design requires that we find a “good” collection of relation schemas.
 - Physical Design – Deciding on the physical layout of the database (like the way to storage data...).

1. Entity Relationship Model

- Entities are described in a database by a set of **attributes**.
- A **relationship** is an association among several entities.

- The set of all entities of the same type and the set of all relationships of the same type are termed **an entity set** and **relationship set**, respectively.
- The overall logical structure (schema) of a database can be expressed graphically by an *entity-relationship (E-R) diagram*.
- There are several ways in which to draw these diagrams. One of the most popular is to use the **Unified Modeling Language (UML)**. In the notation we use, which is based on UML, an E-R diagram is represented as follows:
 - **Entity sets** are represented by a rectangular box with the entity set name in the header and the attributes listed below it.
 - **Relationship sets** are represented by a diamond connecting a pair of related entity sets. The name of the relationship is placed inside the diamond.




Figure 1.3 A sample E-R diagram.

- In addition to entities and relationships, the E-R model represents certain constraints to which the contents of a database must conform. One important constraint is **mapping cardinalities**, which express the number of entities to which another entity can be associated via a relationship set.

2. Normalization

- To determine whether a relation schema is in one of the desirable normal forms, we need additional information about the real-world enterprise that we are modeling with the database. The most common approach is to use **functional dependencies**.
- A bad design may have:
 - Repetition of information
 - Inability to represent certain information

Data Storage and Querying

- The functional components of a database system can be broadly divided into
 - the **storage manager**, and
 - the **query processor components**.

1. Storage Manager

The **storage manager** is the component of a database system that provides the interface between the low-level data stored in the database and the application programs and queries submitted to the system.

- The storage manager is responsible for **storing, retrieving, and updating data** in the database. It contains:
 - **Authorization and integrity manager**
 - **Transaction manager**
 - **File manager**
 - **Buffer manager**
- The storage manager implements several data structures as part of the physical system implementation:
 - **Data files**, which store the database itself.
 - **Data dictionary**, which stores **metadata** about the structure of the database, in particular the schema of the database.
 - **Indices**

2. The Query Processor

- The query processor components include:
 - **DDL interpreter**, which interprets DDL statements and records the definitions in the data dictionary.
 - **DML compiler**, which translates DML statements in a query language into an evaluation plan consisting of low-level instructions that the *query evaluation engine* understands.
 - **Query evaluation engine**, which executes low-level instructions generated by the DML compiler.

A query can usually be translated into any of a number of alternative evaluation plans that all **give the same result**. The DML compiler also performs **query optimization**; that is, it picks the lowest cost evaluation plan from among the alternatives.

Transaction Management

- **ACID** property:
 - The all-or-none requirement is called **atomicity**.
Recovery manager's duty
 - The correctness requirement is called **consistency**.
Concurrency-control manager's duty
 - Each transaction does not feel that other transactions are executing concurrently in the system: **isolation**.

Concurrency-control manager's duty

- The persistence requirement is called **durability**.

Recovery manager's duty

- A **transaction** is a collection of operations that performs a single logical function in a database application.
- The **transaction manager** consists of the **concurrency-control manager** and the **recovery manager**.
- **Failure recovery**: detect system failures and restore the database to the state that existed prior to the occurrence of the failure.